


1 Corpus Sampling

Dolma3 Corpus


Deduplicate

Bin by WebOrganizer:
Topic × Format




Interpretable corpus labels
24 Topics × 24 Formats = 576 Bins

2 Training Data Attribution



Estimate influence of data on output
by measuring similarity of gradients

3 Bin-Level Influence Map



Signed bin influence (z-scores)
All benchmark queries

4 Unlearning Validation

